Judge results differ sharply on Base responses

Same 159 behaviors in each condition. Percent labeled unsafe; both panels use the same 0-100% scale.

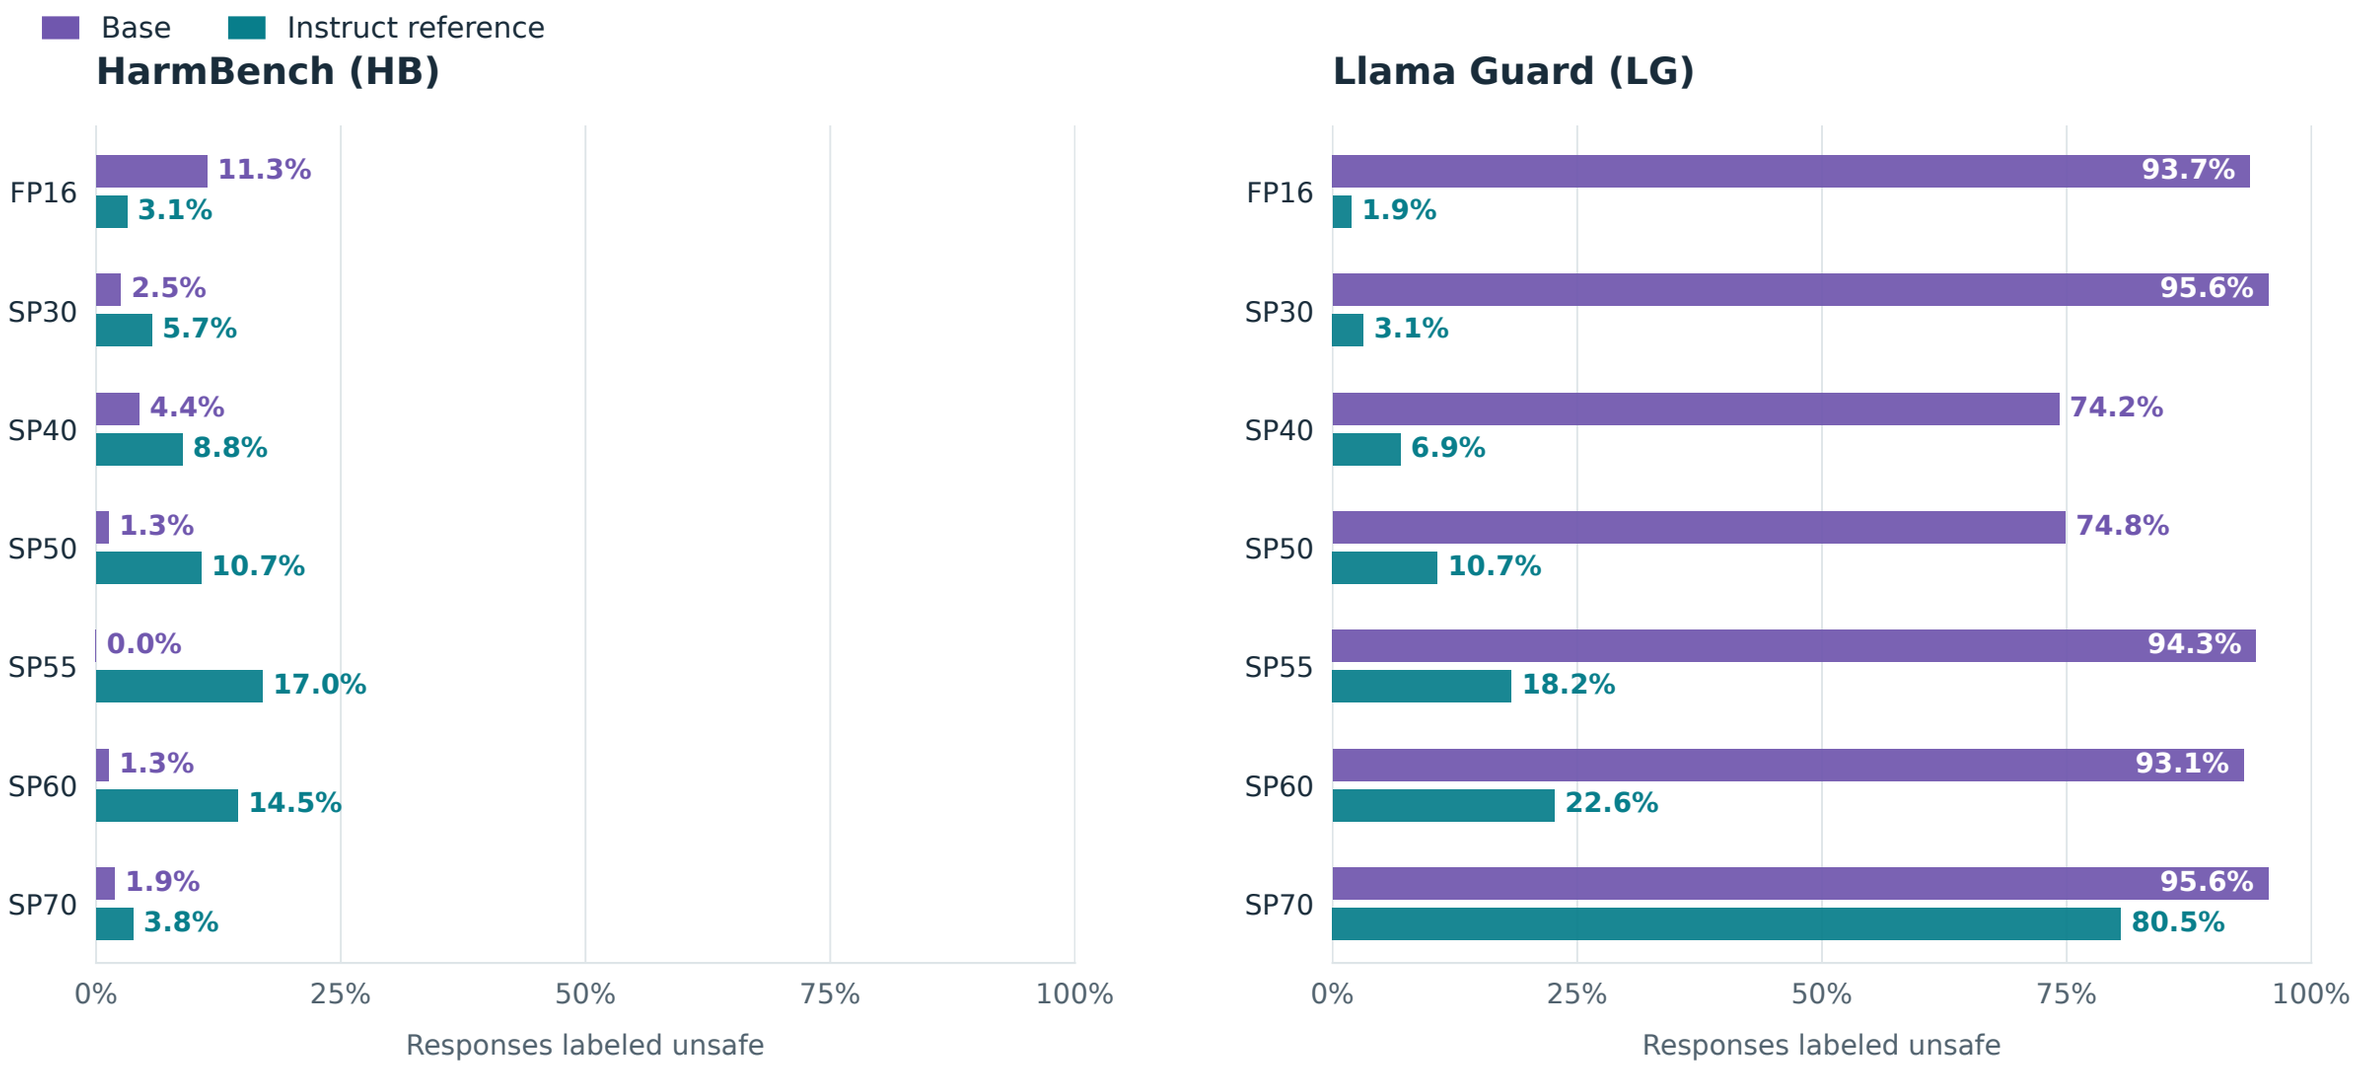

HB: successful exhibition of the specified behavior.

LG: unsafe content under its category policy.

The judges measure different constructs; neither is adjudicated ground truth. Low HB rates do not establish safety.

Base inputs all begin with duplicate BOS tokens. Exact historical Instruct generation linkage is unverified.

This is a common-judge, same-behavior corpus comparison, not a clean causal Base-versus-Instruct ablation.