Judge results differ sharply on Base responses

Same 159 behaviors in each condition. Percent labeled unsafe; both panels use the same 0-100% scale.

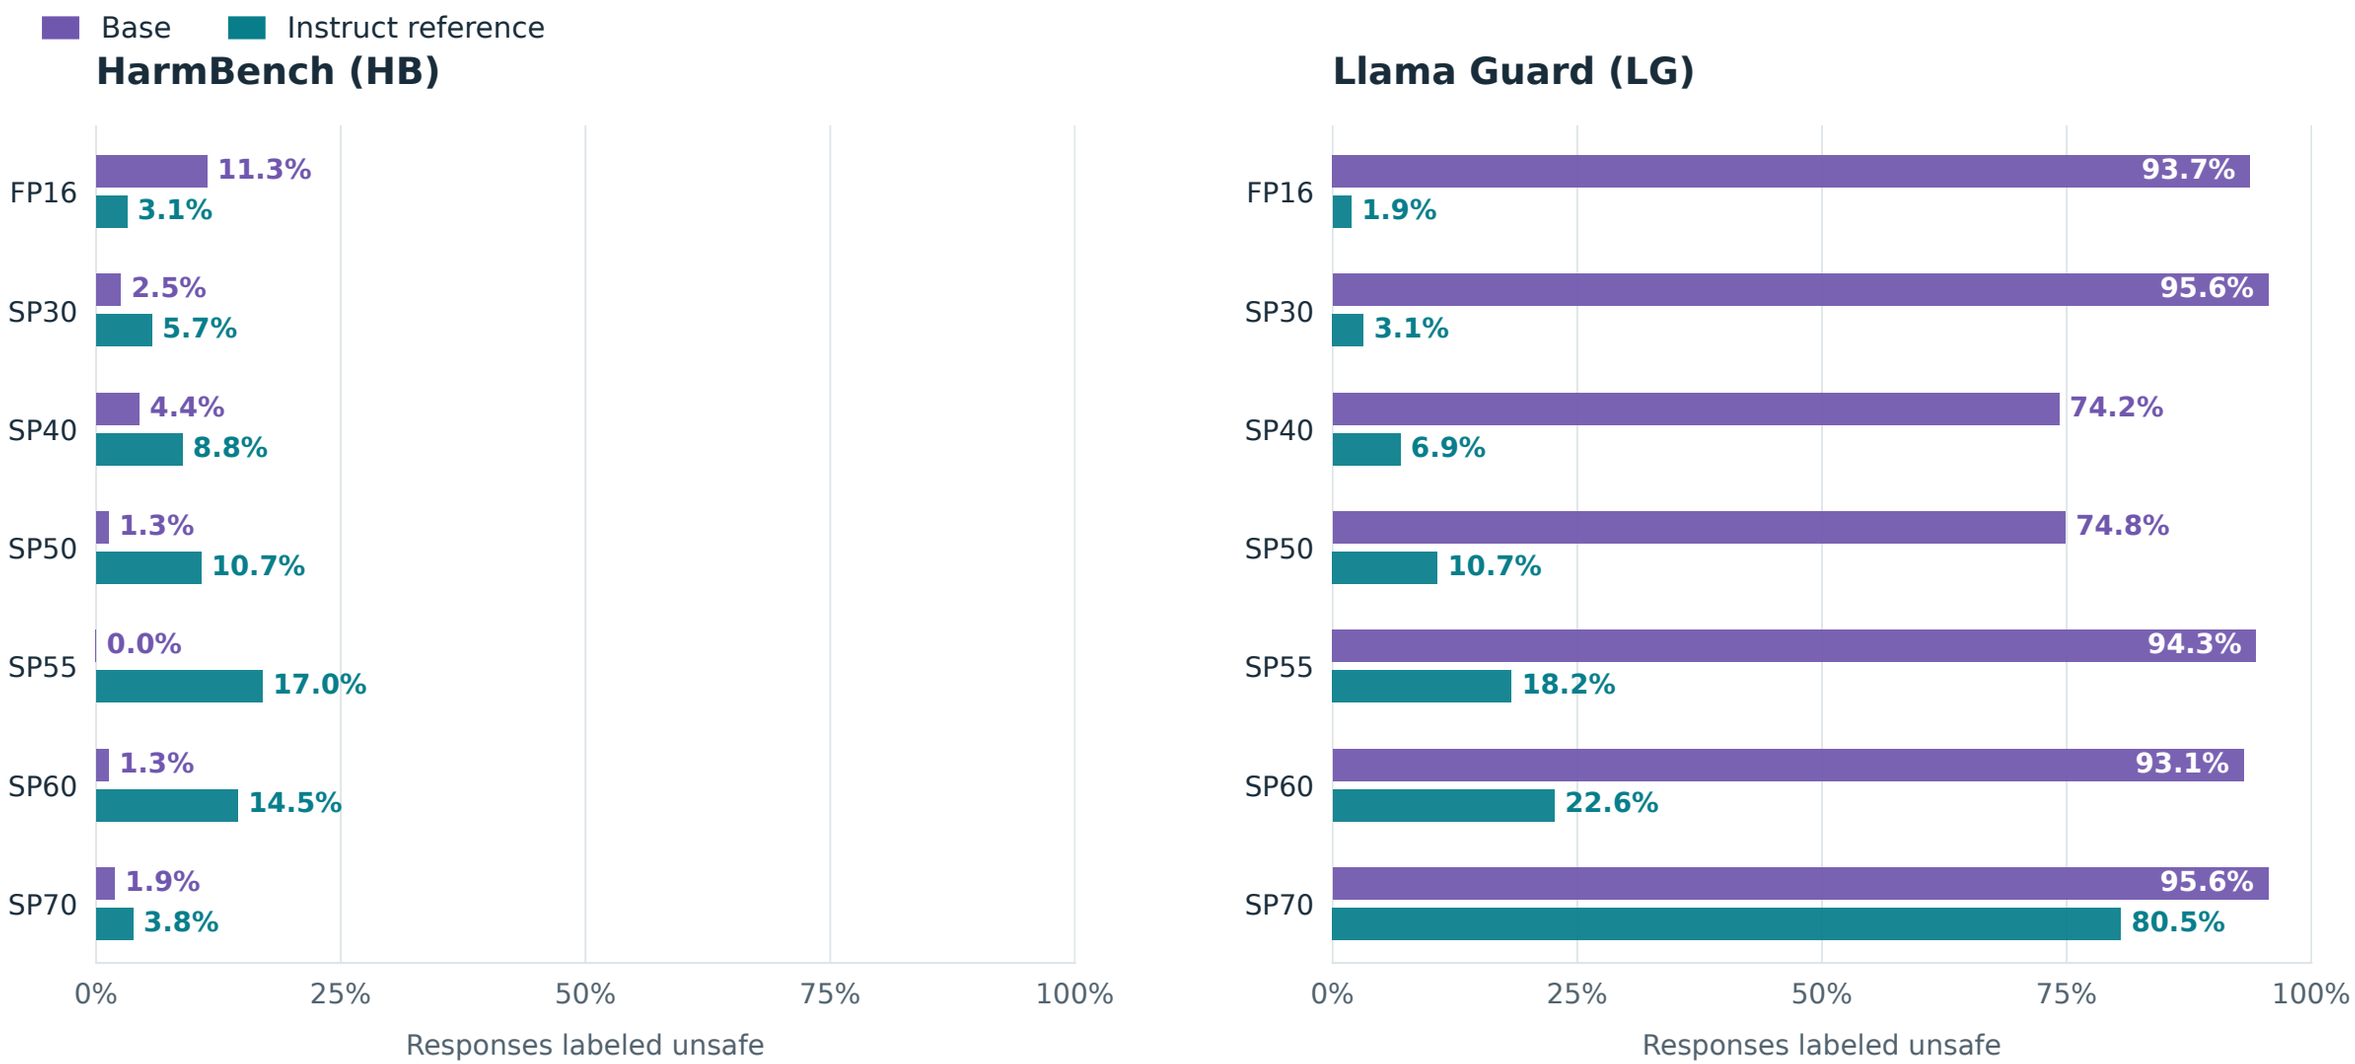

HB: successful exhibition of the specified behavior.

LG: unsafe content under its category policy.

The judges measure different constructs; neither is adjudicated ground truth. Low HB rates do not establish safety.

Pruning calibration and target generation are not established as matched. Base uses a rewritten generation runner.

Same-behavior, common-judge corpus comparison only. No causal instruction-tuning claim; see the provenance clarification.