Counts and interpretation at a glance

Every count has denominator 159. Each cell is Base / Instruct; repetition is a mean score on a 0-1 scale.

Condition	HB unsafe count	LG unsafe count	Judge agreement count	Collapse flag count	Mean word-trigram repetition
FP16	18 / 5	149 / 3	22 / 157	128 / 0	0.822 / 0.006
SP30	4 / 9	152 / 5	11 / 155	148 / 0	0.873 / 0.012
SP40	7 / 14	118 / 11	48 / 154	71 / 0	0.458 / 0.031
SP50	2 / 17	119 / 17	40 / 153	99 / 0	0.589 / 0.048
SP55	0 / 27	150 / 29	9 / 149	156 / 2	0.942 / 0.136
SP60	2 / 23	148 / 36	13 / 138	154 / 10	0.917 / 0.254
SP70	3 / 6	152 / 128	10 / 37	159 / 16	0.964 / 0.225

What the paired tests support

LG labels Base unsafe more often at all seven conditions. HB is higher for Base at FP16 and lower at SP50, SP55 and SP60.
HB differences at SP30, SP40 and SP70 do not survive the 14-test Holm correction (exact two-sided McNemar).
At SP70, the HB difference is -1.89 percentage points (Holm p = 0.534): this does not show that Base is safer.

Interpretation boundary

Verified: 159 shared behaviors, 1,113 response pairs, 2,226 paired judge labels; reference labels and hashes match the original Oct 29 / Oct 30 judged records. Matching judge settings do not make generation pipelines identical.
All Base inputs start with duplicate BOS. Historical Instruct prompt/code/tokenizer/runtime and checkpoint linkage are not established here. Neither judge resolves the other judge's labels; no causal instruction-tuning claim follows.

Collapse proxy: >=50 words with unique-word ratio <0.10 or trigram repetition >=0.80; or >=20 words with dominance >=0.30.
The same 159 behaviors recur across conditions. The 1,113 response pairs are not 1,113 independent behaviors.
Public reproducibility: aggregate metrics to figures and tables only. No original prompts, responses or individual judge records.